

Q10 Perform compression on mnist dataset using auto encoder.

Aim: to compress MNIST hand written digit image into a lower-dimensional representation using a autoencoder.

Algorithm

- 1, load and normalize the MNIST dataset
- 2, built an autoencoder neural network
- 3, train the autoencoders to minimize the difference between original and reconstructed image
- 4, use the encoder to compress the image
- 5, use the decoder to reconstruct the image

pseudo code:

load MNIST dataset

Normalize images to range [0, 1]

Define encoder network:

input : 28x28 image

flatten input

dense layers reducing dimension \rightarrow latent space


Define decoder network:

input : latent vector

dense layers expanding dimension

reshape output to 28x28 image

Autoencoder architecture



reconstructed
image

combine encoders with loss = mean-squared error

train autoencoder on MNIST training images

for a given image:

compressed = encoder (image)

reconstructed = decoder (compressed)

Result:

+ typical reconstruction loss reduces significantly after ~~steps~~ training, showing the auto encoder successfully

~~eff. 128~~

~~combine encoders and decoder to form~~

out put :

Epoch : [1/10] loss : 0.9496

Epoch : [2/10] loss : 0.9254

Epoch : [3/10] loss : 0.9254

Epoch : [4/10] loss : 0.9254

Epoch : [5/10] loss : 0.9254

Epoch : [6/10] loss : 0.9254






Epoch (7/10) loss: 0.9254

Epoch (8/10) loss: 0.9254





Epoch (9/10) loss: 0.9254

Epoch (10/10) loss: 0.9254

Auto coder training loss



original images



Reconstructed images

```
import tensorflow as tf
from tensorflow.keras import layers, Model
import matplotlib.pyplot as plt
import numpy as np

(x_train, _), (x_test, _) = tf.keras.datasets.mnist.load_data()
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0
x_train = x_train.reshape(-1, 784)
x_test = x_test.reshape(-1, 784)

latent_dim = 2
inputs = tf.keras.Input(shape=(784,), name="encoder_input")
x = layers.Dense(512, activation="relu")(inputs)
z_mean = layers.Dense(latent_dim, name="z_mean")(x)
z_log_var = layers.Dense(latent_dim, name="z_log_var")(x)

def sampling(args):
    z_mean, z_log_var = args
    epsilon = tf.random.normal(shape=tf.shape(z_mean))
    return z_mean + tf.exp(0.5 * z_log_var) * epsilon

z = layers.Lambda(sampling, name="z")([z_mean, z_log_var])
encoder = Model(inputs, [z_mean, z_log_var, z], name="encoder")
encoder.summary()

latent_inputs = tf.keras.Input(shape=(latent_dim,), name="z_sampling")
x = layers.Dense(512, activation="relu")(latent_inputs)
outputs = layers.Dense(784, activation="sigmoid")(x)
decoder = Model(latent_inputs, outputs, name="decoder")
decoder.summary()

class VAE(Model):
    def __init__(self, encoder, decoder, **kwargs):
        super(VAE, self).__init__(**kwargs)
        self.encoder = encoder
        self.decoder = decoder

    def call(self, inputs):
        z_mean, z_log_var, z = self.encoder(inputs)
        return self.decoder(z)

    def train_step(self, data):
        if isinstance(data, tuple):
            data = data[0]
        with tf.GradientTape() as tape:
            z_mean, z_log_var, z = self.encoder(data)
            reconstruction = self.decoder(z)
            # ✅ Fixed reconstruction loss
            reconstruction_loss = tf.reduce_mean(
                tf.keras.losses.binary_crossentropy(data, reconstruction) * 784
            )
            kl_loss = -0.5 * tf.reduce_mean(
                1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var)
            )
            total_loss = reconstruction_loss + kl_loss
        grads = tape.gradient(total_loss, self.trainable_variables)
        self.optimizer.apply_gradients(zip(grads, self.trainable_variables))
        return {"loss": total_loss}

    def test_step(self, data):
        pass
```

