

## Ex 11 → Variational Autoencoder (VAE)

7/10

Aim: to build a variational autoencoder that can learn compressed representations of data and generate new similar samples

Algorithm:

- 1, load and normalize the dataset
- 2, pass data through an encoder to get two outputs
- 3, use the reparameterization trick.
- 4, pass  $z$  through a decoder to reconstruct the input
- 5, compute total loss = reconstruction loss + KL divergence
- 6, update network weights using a optimizer
- 7, Repeat for multiple epochs.

pseudo code:

initialize encoder and decoder

for each epoch:

    for each batch of data

$\mu, \log\_var = \text{encoder}(x)$


$z = \mu + \text{exp}(\log\_var) * \text{random\_noise}()$

$\hat{x} = \text{decoder}(z)$

$\text{loss} = \text{reconstruction\_loss}(x, \hat{x})$

            + KL-divergence( $\mu, \log\_var$ )

variational auto encoder architecture.



## update . weights ( loss )

### Observation :

- \* the loss decreases with training.
- \* re constructed images become clearer after several epochs
- \* model learns a smooth latent space - similar datapoint cluster together


### Result :

- ~~\* the VAE successfully compresses and reconstructs input data.~~

~~extra info~~

Epoch [1/10] loss 191.10 EPOCH [6/10] loss : 157.21  
 EPOCH [2/10] loss 161.28 EPOCH [7/10] loss : 156.11  
 EPOCH [3/10] loss 162.28 EPOCH [8/10] loss : 155.11  
 EPOCH [4/10] loss 162.21 EPOCH [9/10] loss : 154.11  
 EPOCH [5/10] loss 160.63 EPOCH [10/10] loss : 154.11

VAE training loss



```
from tensorflow.keras.layers import Input, Dense
import matplotlib.pyplot as plt
import numpy as np
(x_train, _), (x_test, _) = mnist.load_data()          # labels not needed
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0

x_train = x_train.reshape(len(x_train), 784)
x_test = x_test.reshape(len(x_test), 784)

input_img = Input(shape=(784,))
encoded = Dense(128, activation='relu')(input_img)
encoded = Dense(64, activation='relu')(encoded)
encoded = Dense(32, activation='relu')(encoded)          # compressed layer

decoded = Dense(64, activation='relu')(encoded)
decoded = Dense(128, activation='relu')(decoded)
decoded = Dense(784, activation='sigmoid')(decoded)     # reconstruct pixels

autoencoder = Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='mse', metrics=['accuracy'])

history = autoencoder.fit(
    x_train, x_train,
    epochs=10,
    batch_size=256,
    shuffle=True,
    validation_data=(x_test, x_test)
)

plt.figure(figsize=(6,4))
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title("Autoencoder Training Loss")
plt.xlabel("Epochs")
plt.ylabel("MSE Loss")
plt.legend()
plt.show()

decoded_imgs = autoencoder.predict(x_test)
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    plt.title("Original")
    plt.axis("off")

    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
    plt.title("Reconstructed")
    plt.axis("off")

plt.show()
```

Model: "encoder"

Layer (type)	Output Shape	Param #	Connected to
encoder_input (InputLayer)	(None, 784)	8	-
dense_47 (Dense)	(None, 512)	401,920	encoder_input[0]..
z_mean (Dense)	(None, 2)	1,026	dense_47[0][0]
z_log_var (Dense)	(None, 2)	1,026	dense_47[0][0]
z (Lambda)	(None, 2)	0	z_mean[0][0], z_log_var[0][0]

Total params: 403,972 (1.54 MB)

Trainable params: 403,972 (1.54 MB)

Non-trainable params: 0 (0.00 B)

Model: "decoder"

Layer (type)	Output Shape	Param #
z_sampling (InputLayer)	(None, 2)	0
dense_48 (Dense)	(None, 512)	1,536
dense_49 (Dense)	(None, 784)	402,192

Total params: 403,728 (1.54 MB)

Trainable params: 403,728 (1.54 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

235/235 - 14s 58ms/step - loss: 210.1661 - val\_loss: 0.0000e+00

Epoch 2/10

235/235 - 9s 37ms/step - loss: 171.5112 - val\_loss: 0.0000e+00

Epoch 3/10

235/235 - 10s 42ms/step - loss: 164.4683 - val\_loss: 0.0000e+00

Epoch 4/10

235/235 - 10s 42ms/step - loss: 161.6138 - val\_loss: 0.0000e+00

Epoch 5/10

235/235 - 10s 40ms/step - loss: 159.8511 - val\_loss: 0.0000e+00

Epoch 6/10

235/235 - 9s 39ms/step - loss: 158.4258 - val\_loss: 0.0000e+00

Epoch 7/10

235/235 - 11s 44ms/step - loss: 157.1982 - val\_loss: 0.0000e+00

Epoch 8/10

235/235 - 10s 43ms/step - loss: 156.0260 - val\_loss: 0.0000e+00

Epoch 9/10

235/235 - 9s 40ms/step - loss: 154.9950 - val\_loss: 0.0000e+00

Epoch 10/10

235/235 - 9s 39ms/step - loss: 154.0985 - val\_loss: 0.0000e+00

